

Imports

```
In [6]: import collections
import itertools
import os
import random
from pathlib import Path
from typing import Dict, Iterable, List, Optional, Tuple

import albumentations as A
import cv2
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import torch
import torchmetrics
from albumentations.augmentations import transforms as AT
from albumentations.augmentations.dropout.coarse_dropout import CoarseDropout
from albumentations.pytorch import ToTensorV2
from madgrad import MADGRAD
from sklearn.model_selection import train_test_split
from torch import nn
from torch.utils import data
from torchvision import transforms
from tqdm import tqdm

cv2.setNumThreads(4)
cv2.ocl.setUseOpenCL(True)

def set_seed(seed):
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False
    np.random.seed(seed)
    random.seed(seed)
    os.environ["PYTHONHASHSEED"] = str(seed)

set_seed(33)
```

Classes

```
In [2]: class CharEncoder:
    """Encode and decode chars

    Args:
        alphabet (str): string of all chars
    """
    def __init__(self, alphabet: str):
        alphabet = alphabet.lower()
        self._alphabet = alphabet

        self._encode_dict = {}
        for i, char in enumerate(alphabet):
            self._encode_dict[char] = i + 1

        self._decode_dict = {v: k for k, v in self._encode_dict.items()}
        self._decode_dict[0] = ""

    def encode(self, text: Iterable[str]) -> torch.Tensor:
        """Encode chars

        Args:
            text (Iterable[str]): set of strings to encode

        Returns:
            torch.Tensor: flat Tensor with encoded chars
        """
        text = "".join(text).lower()
        text = list(map(self._encode_dict.get, text))
        return torch.IntTensor(text)

    def decode(self, x: torch.Tensor, length: int) -> np.ndarray:
        """Decode numbers

        Args:
            x (torch.Tensor): flat Tensor with encoded chars
```

```

        length (int): length of one word

Returns:
    np.ndarray: ndarray with decoded numbers
    """
    x = np.array(list(map(lambda x: self._decode_dict.get(x.item()), x)))
    x = np.split(x, len(x) // length)
    x = np.apply_along_axis(
        lambda k: "".join(c[0] for c in itertools.groupby(k)), 1, x
    )
    return np.array(x)

@property
def alphabet(self):
    return self._alphabet

```

```

In [5]: class MyDataset(data.Dataset):
        """Dataset for CCPD2019

        Args:
            files (Iterable[Path]): paths to all images
            transform (A.Compose): transforms from albumentations
            alphabet (bool, optional): Create an alphabet from filenames or not. Defaults to True.
        """
        def __init__(
            self, files: Iterable[Path], transform: A.Compose, alphabet: bool = True
        ) -> None:
            super().__init__()
            self._files = files
            self._labels = []
            self._alphabet = alphabet
            self.transform = transform
            if alphabet is True:
                self._alphabet = set()
            for file in files:
                label = file.stem.split("-")[1].lower()
                self._labels.append(label)

                if alphabet is True:
                    self._alphabet |= set(label)

        @property
        def alphabet(self) -> str:
            if not self._alphabet:
                raise AttributeError()
            return self._alphabet

        def load_sample(self, filepath: Path) -> np.ndarray:
            """Load image via cv2

            Args:
                filepath (Path): path to image
            Returns:
                np.ndarray: loaded image
            """
            image = cv2.imdecode(np.fromfile(filepath, dtype=np.uint8), cv2.IMREAD_COLOR)
            image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
            return image

        def __len__(self) -> int:
            return len(self._files)

        def __getitem__(self, index: int) -> Tuple[torch.Tensor, str]:
            x = self.load_sample(self._files[index])
            x = self.transform(image=x)["image"] / 256
            y = self._labels[index]
            return x, y

```

```

In [6]: class Accuracy:
        """Accuracy metric
        """
        def __call__(self, y_pred: np.ndarray, y_true: np.ndarray) -> float:
            """Calculate accuracy

            Args:
                y_pred (np.ndarray): predicted labels
                y_true (np.ndarray): true labels

            Returns:
                float: accuracy
            """
            if y_pred.shape != y_true.shape:

```

```
        raise AttributeError
    return np.sum(y_true==y_pred)/len(y_true)
```

```
In [7]: class BiLSTM(nn.Module):
        """Bidirectional LSTM

        Args:
            n_channels (int): number of expected features in the input x
            n_hidden (int): number of features in the hidden state h
            n_out (int): number of out features
        """
        def __init__(self, n_channels: int, n_hidden: int, n_out: int):

            super(BiLSTM, self).__init__()
            self.rnn = nn.LSTM(n_channels, n_hidden, bidirectional=True, batch_first=True)
            self.linear = nn.Linear(n_hidden * 2, n_out)

        def forward(self, input: torch.Tensor):
            recurrent, _ = self.rnn(input)
            output = self.linear(recurrent)
            return output

class FCNN(nn.Module):
    """Fully-convolutional CNN

    Args:
        n_out (int): number of out features
    """
    def __init__(self, n_out: int):
        super(FCNN, self).__init__()
        #in 3x32x128
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, 3, 1, 1),
            nn.ReLU(True),
            nn.MaxPool2d((2, 1), (2, 1)),
        )
        #out 64x16x128

        self.conv2 = nn.Sequential(
            nn.Conv2d(64, 128, 3, 1, 1),
            nn.ReLU(True),
            nn.MaxPool2d(2, 2),
        )
        #out 128x8x64

        self.conv3 = nn.Sequential(
            nn.Conv2d(128, 256, 3, 1, 1),
            nn.BatchNorm2d(256),
            nn.ReLU(True),
            nn.Conv2d(256, 256, 3, 1, 1),
            nn.ReLU(True),
            nn.MaxPool2d((2, 1), (2, 1)),
        )
        #out 256x4x64

        self.conv4 = nn.Sequential(
            nn.Conv2d(256, 512, 3, 1, 1),
            nn.BatchNorm2d(512),
            nn.ReLU(True),
            nn.Conv2d(512, 512, 3, 1, 1),
            nn.ReLU(True),
            nn.MaxPool2d(2, 2),
        )
        #out 512x2x32

        self.conv5 = nn.Sequential(
            nn.Conv2d(512, 512, 3, 1, 1),
            nn.BatchNorm2d(512),
            nn.ReLU(True),
            nn.MaxPool2d((2, 1), (2, 1)),
        )
        #out 512x1x32

        self.fc = nn.Sequential(
            nn.Flatten(2),
            nn.Linear(32, 256),
            nn.ReLU(True),
            nn.Linear(256, n_out)
        )

    def forward(self, input: torch.Tensor):
        x = self.conv1(input)
        x = self.conv2(x)
        x = self.conv3(x)
```

```

        x = self.conv4(x)
        x = self.conv5(x)
        x = self.fc(x)
        return x

class Model(nn.Module):
    """Model for car license plate recognition

    Args:
        alphabet_length (_type_): length of alphabet
        sequence_length (int, optional): FCNN output width. Defaults to 16.
        lstm_layers (int, optional): number of LSTM hidden layers. Defaults to 256.
    """
    def __init__(self, alphabet_length, sequence_length=16, lstm_layers=256):
        super(Model, self).__init__()

        self.alphabet_length = alphabet_length
        self.sequence_length = sequence_length
        self.cnn = FCNN(sequence_length)
        self.rnn = nn.Sequential(
            BiLSTM(512, lstm_layers, lstm_layers),
            BiLSTM(lstm_layers, lstm_layers, alphabet_length))

    def forward(self, x):
        x = self.cnn(x).permute(0, 2, 1) # b, w, c
        output = self.rnn(x)
        return output

```

```

In [8]: class OCR:
        """OCR Model trainer

        Args:
            model (nn.Module): model to train
            optimizer (torch.optim.Optimizer): optimizer
            criterion (nn.CTCLoss): criterion
            target_length (int): length of target
            alphabet (str): string of all possible target characters
            gpu_transform (Optional[transforms.Compose], optional):
                transforms that produced by the GPU. Defaults to None.
            scheduler (Optional[torch.optim.lr_scheduler._LRScheduler], optional):
                scheduler. Defaults to None.
            device (str, optional): device. Defaults to "cuda".
        """

        def __init__(
            self,
            model: nn.Module,
            optimizer: torch.optim.Optimizer,
            criterion: nn.CTCLoss,
            target_length: int,
            alphabet: str,
            gpu_transform: Optional[transforms.Compose] = None,
            scheduler: Optional[torch.optim.lr_scheduler._LRScheduler] = None,
            device="cuda",
        ) -> None:
            self.model = model
            self.optimizer = optimizer
            self.criterion = criterion
            self.target_length = target_length
            self.alphabet = alphabet
            self.gpu_transform = gpu_transform
            self.scheduler = scheduler
            self.device = device
            self.label_encoder = CharEncoder(alphabet)

            self._best_score = -1
            self._best_weights = None

            if len(alphabet) + 1 != model.alphabet_length:
                raise ValueError(
                    "The size of the model target does not match the specified one"
                )

        def _count_phase_metrics(self, phase_metrics: Dict, batch_metrics: Dict):
            """Add batcg metrics to phase metrics

            Args:
                phase_metrics (Dict): dict of phase metrics
                batch_metrics (Dict): dict of batch metrics
            """
            for metric in batch_metrics:

```

```

        phase_metrics[metric] += batch_metrics[metric]

def _batch_handler(
    self,
    x_batch: torch.Tensor,
    y_batch: torch.Tensor,
    y_decoded: np.ndarray,
    metrics: Iterable,
) -> Tuple[torch.Tensor, Dict]:
    """Handle batch

    Args:
        x_batch (torch.Tensor): images batch
        y_batch (torch.Tensor): encoded targets batch
        y_decoded (np.ndarray): decoded targets batch
        metrics (Iterable): metrics to calculate

    Returns:
        Tuple[torch.Tensor, Dict]: loss and dict of metrics
    """
    outputs, preds = self.predict(x_batch)
    loss = self.criterion(
        outputs.permute(1, 0, 2).log_softmax(-1), # l, b, c
        y_batch,
        torch.IntTensor([self.model.sequence_length] * len(y_batch)),
        torch.IntTensor([self.target_length] * len(y_batch)),
    )
    counted_metrics = {
        metric.__class__.__name__: metric(preds, y_decoded) * len(y_batch)
        for metric in metrics
    }
    return loss, counted_metrics

def _phase_handler(
    self,
    phase: str,
    dataloader: data.DataLoader,
    metrics: Iterable,
    load_best_weights: Optional[str] = None,
) -> Tuple[float, Dict]:
    """Handle phase

    Args:
        phase (str): phase name
        dataloader (data.DataLoader): dataloader
        metrics (Iterable): metrics to calculate
        load_best_weights (Optional[str], optional):
            by which metric best weights of model will be loaded at the end of training.
            If None, weights will not be loaded. Defaults to None.

    Returns:
        Tuple[float, Dict]: loss and dict of metrics
    """
    phase_loss = 0
    phase_metrics = collections.defaultdict(lambda: 0)
    for x_batch, y_batch in tqdm(dataloader, desc=f"Phase {phase}"):
        x_batch = x_batch.to(self.device)

        y_decoded = np.array(y_batch)
        y_batch = self.label_encoder.encode(y_batch)
        y_batch = y_batch.to(self.device).view(-1, self.target_length)

        if phase == "train":
            if self.gpu_transform is not None:
                x_batch = self.gpu_transform(x_batch)
            self.optimizer.zero_grad()

        batch_loss, batch_metrics = self._batch_handler(
            x_batch, y_batch, y_decoded, metrics
        )

        if phase == "train":
            batch_loss.backward()
            self.optimizer.step()

        phase_loss += batch_loss.item() * len(y_batch)
        self._count_phase_metrics(phase_metrics, batch_metrics)

    phase_loss /= len(dataloader.dataset)
    phase_metrics = {
        metric: value / len(dataloader.dataset)
        for metric, value in phase_metrics.items()
    }

```

```

    if (
        phase == "val"
        and load_best_weights
        and phase_metrics[load_best_weights] > self._best_score
    ):
        self._best_score = phase_metrics[load_best_weights]
        self._best_weights = self.model.state_dict()

    tqdm.write(f"\tLoss: {phase_loss:0.5f}")
    for metric, value in phase_metrics.items():
        tqdm.write(f"\t{metric}: {value:0.5f}")
    print()
    return phase_loss, dict(phase_metrics)

def fit(
    self,
    train: data.DataLoader,
    val: data.DataLoader,
    num_epochs: int,
    metrics: Iterable,
    load_best_weights: Optional[str] = "Accuracy",
    scheduler_metric: Optional[str] = None,
) -> Tuple[List, List]:
    """fit model

    Args:
        train (data.DataLoader): train dataloader
        val (data.DataLoader): val dataloader
        num_epochs (int): number of epochs
        metrics (Iterable): metrics to calculate
        load_best_weights (Optional[str], optional): by which metric
            best weights of the model will be loaded at the end of training.
            If None, weights will not be loaded. Defaults to "Accuracy"
        scheduler_metric (Optional[str], optional):
            needed if scheduler requires metric for step. Defaults to None.

    Returns:
        Tuple[List, List]: train losses and val losses
    """
    if scheduler_metric is not None and self.scheduler is None:
        raise ValueError("Scheduler is not specified")
    if load_best_weights not in metrics:
        raise ValueError("{load_best_weights} is not specified in metrics")

    train_losses = []
    val_losses = []

    for epoch in range(num_epochs):
        tqdm.write(f"Epoch {epoch+1:03d}")

        self.model.train()
        train_loss, train_metrics = self._phase_handler("train", train, metrics)

        self.model.eval()
        with torch.no_grad():
            val_loss, val_metrics = self._phase_handler(
                "val", val, metrics, load_best_weights
            )

        train_losses.append(train_loss)
        val_losses.append(val_loss)

        if self.scheduler:
            tqdm.write(f"\nLR: {self.optimizer.param_groups[0]['lr']}")
            if scheduler_metric:
                self.scheduler.step(val_metrics[scheduler_metric])
            else:
                self.scheduler.step()

        tqdm.write("-" * 40)

    if load_best_weights:
        self.model.load_state_dict(self._best_weights)
    self._best_weights = None

    return train_losses, val_losses

@torch.no_grad()
def score_test(
    self, dataloader: data.DataLoader, metrics: Iterable
) -> Tuple[float, Dict]:
    """Calculate metrics on test data

```

```

    Args:
        dataloader (data.DataLoader): test dataloader
        metrics (Iterable): metrics to calculate

    Returns:
        Tuple[float, Dict]: loss and dict of metrics
    """
    self.model.eval()
    test_loss, test_metrics = self._phase_handler("test", dataloader, metrics)
    return test_loss, test_metrics

@torch.no_grad()
def get_test_scores(
    self, dataset: data.Dataset, metric: torchmetrics.Metric, device: str = "cpu"
) -> List:
    """Get score for each image in test

    Args:
        dataset (data.Dataset): test dataset
        metric (torchmetrics.Metric): metric to calculate
        device (str, optional): device. Defaults to "cpu".

    Returns:
        List: list of scores
    """
    self.model.to(device)
    self.model.eval()

    metric_list = []
    for x, y in tqdm(dataset):
        x = x.to(device)
        outputs, preds = self.predict(x)
        metric_list.append(
            metric(
                preds,
                [
                    y,
                ],
            ).item()
        )

    self.model.to(self.device)
    return metric_list

def predict(self, x: torch.Tensor) -> Tuple[torch.Tensor, np.ndarray]:
    """predict

    Args:
        x (torch.Tensor): one or more images

    Returns:
        Tuple[torch.Tensor, np.ndarray]: raw output, decoded output
    """
    if x.dim() == 3:
        x = x.unsqueeze(0)
    outputs = self.model(x)
    preds = self.label_encoder.decode(
        torch.argmax(outputs, 2).flatten(), self.model.sequence_length
    )
    return outputs, preds

def state_dict(self) -> Dict:
    """Get state dict of:
    model, optimizer, scheduler, target_length and alphabet

    Returns:
        Dict: state dict
    """
    state_dict = {
        "model": self.model.state_dict(),
        "optimizer": self.optimizer.state_dict(),
        "scheduler": self.scheduler.state_dict(),
        "target_length": self.target_length,
        "alphabet": self.alphabet,
    }
    return state_dict

def load_state_dict(
    self, state_dict: Dict, load_optimizer: bool = True, load_scheduler: bool = True
) -> None:
    """load state dict

```

```

    Args:
        state_dict (Dict): state dict
        load_optimizer (bool, optional): load the optimizer or not. Defaults to True.
        load_scheduler (bool, optional): load the scheduler or not. Defaults to True.
    """
    self.model.load_state_dict(state_dict["model"])
    self.target_length = state_dict["target_length"]
    self.alphabet = state_dict["alphabet"]

    if load_optimizer:
        self.optimizer.load_state_dict(state_dict["optimizer"])
    if load_scheduler:
        self.scheduler.load_state_dict(state_dict["scheduler"])

```

Model Training

```

In [ ]: TRAIN_DIR = Path("data/train/")
        TEST_DIR = Path("data/test/")
        WORD_LENGTH = 7

```

```

In [9]: train_val_files = list(TRAIN_DIR.rglob("*.jpg"))
        train_files, val_files = train_test_split(train_val_files, test_size=0.1)
        test_files = list(TEST_DIR.rglob("*.jpg"))

```

```

In [10]: train_transform = A.Compose([
        A.Resize(32, 128),
        A.Emboss((0.2, 1), p=0.5),
        A.PixelDropout(0.1, drop_value = (200,200,200), p=0.3),
        CoarseDropout(max_holes=5, min_holes=2, min_height=3, min_width=3, fill_value=(200,200,200), p=0.2),
        A.GaussNoise((50, 150), p=0.5),
        AT.ColorJitter((0.5, 1.1), hue=0.2, p=0.7),
        ToTensorV2(),

    ])

    val_transform = A.Compose([
        A.Resize(32, 128),
        ToTensorV2(),

    ])

    gpu_transform = transforms.Compose([
        transforms.RandomPerspective(0.2, p=0.5),

    ])

```

```

In [11]: train_dataset = MyDataset(train_files, train_transform)
        val_dataset = MyDataset(val_files, val_transform)
        test_dataset = MyDataset(test_files, val_transform)

        alphabet = "".join(
            train_dataset.alphabet | val_dataset.alphabet | test_dataset.alphabet
        )

```

```

In [12]: train = data.DataLoader(train_dataset, batch_size=400, shuffle=True, drop_last=True)
        val = data.DataLoader(val_dataset, batch_size=400, drop_last=True)
        test = data.DataLoader(test_dataset, batch_size=400)

```

```


In [13]: plt.imshow(gpu_transform(train_dataset[1223][0]).permute(1,2,0))

```

```

Out[13]: <matplotlib.image.AxesImage at 0x20b87faa940>

```



```

In [14]: device = "cuda"
        model = Model(len(alphabet) + 1).to(device)
        loss = nn.CTCLoss()
        optimizer = MADGRAD(model.parameters(), lr=1e-3)
        scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
            optimizer,
            mode="max",
            factor=0.1,
            patience=1,
            threshold=0.0005,
            threshold_mode="abs",

        )

```



```
In [15]: ocr = OCR(  
    model,  
    optimizer,  
    loss,  
    WORD_LENGTH,  
    alphabet,  
    gpu_transform,  
    scheduler=scheduler,  
    device=device  
)
```

```
In [16]: n_epochs = 15  
metrics = [Accuracy(), torchmetrics.CharErrorRate()]  
train_losses, val_losses = ocr.fit(train, val, n_epochs, metrics, scheduler_metric='Accuracy')
```

Epoch 001

```
Phase train: 100%|██████████| 449/449 [08:50<00:00, 1.18s/it]  
    Loss: 1.85636  
    Accuracy: 0.16110  
    CharErrorRate: 0.52306
```

```
Phase val: 100%|██████████| 49/49 [00:26<00:00, 1.87it/s]  
    Loss: 0.07960  
    Accuracy: 0.88764  
    CharErrorRate: 0.01428
```

LR: 0.001

Epoch 002

```
Phase train: 100%|██████████| 449/449 [09:28<00:00, 1.27s/it]  
    Loss: 0.08450  
    Accuracy: 0.88806  
    CharErrorRate: 0.01975
```

```
Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.91it/s]  
    Loss: 0.01329  
    Accuracy: 0.95865  
    CharErrorRate: 0.00319
```

LR: 0.001

Epoch 003

```
Phase train: 100%|██████████| 449/449 [09:25<00:00, 1.26s/it]  
    Loss: 0.02405  
    Accuracy: 0.95763  
    CharErrorRate: 0.00629
```

```
Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.95it/s]  
    Loss: 0.00746  
    Accuracy: 0.96880  
    CharErrorRate: 0.00169
```

LR: 0.001

Epoch 004

```
Phase train: 100%|██████████| 449/449 [09:37<00:00, 1.29s/it]  
    Loss: 0.01495  
    Accuracy: 0.97314  
    CharErrorRate: 0.00383
```

```
Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.90it/s]  
    Loss: 0.00551  
    Accuracy: 0.97120  
    CharErrorRate: 0.00133
```

LR: 0.001

Epoch 005

```
Phase train: 100%|██████████| 449/449 [08:49<00:00, 1.18s/it]  
    Loss: 0.01491  
    Accuracy: 0.97338  
    CharErrorRate: 0.00389
```

```
Phase val: 100%|██████████| 49/49 [00:24<00:00, 2.00it/s]
```

Loss: 0.00545
Accuracy: 0.97195
CharErrorRate: 0.00129

LR: 0.001

Epoch 006

Phase train: 100%|██████████| 449/449 [08:59<00:00, 1.20s/it]

Loss: 0.00962
Accuracy: 0.98185
CharErrorRate: 0.00249

Phase val: 100%|██████████| 49/49 [00:24<00:00, 1.97it/s]

Loss: 0.00294
Accuracy: 0.97600
CharErrorRate: 0.00064

LR: 0.001

Epoch 007

Phase train: 100%|██████████| 449/449 [09:19<00:00, 1.25s/it]

Loss: 0.00871
Accuracy: 0.98283
CharErrorRate: 0.00229

Phase val: 100%|██████████| 49/49 [00:26<00:00, 1.82it/s]

Loss: 0.00335
Accuracy: 0.97500
CharErrorRate: 0.00079

LR: 0.001

Epoch 008

Phase train: 100%|██████████| 449/449 [09:21<00:00, 1.25s/it]

Loss: 0.00707
Accuracy: 0.98574
CharErrorRate: 0.00188

Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.94it/s]

Loss: 0.00246
Accuracy: 0.97650
CharErrorRate: 0.00056

LR: 0.001

Epoch 009

Phase train: 100%|██████████| 449/449 [09:24<00:00, 1.26s/it]

Loss: 0.00581
Accuracy: 0.98760
CharErrorRate: 0.00158

Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.92it/s]

Loss: 0.00187
Accuracy: 0.97755
CharErrorRate: 0.00042

LR: 0.001

Epoch 010

Phase train: 100%|██████████| 449/449 [09:21<00:00, 1.25s/it]

Loss: 0.00516
Accuracy: 0.98659
CharErrorRate: 0.00170

Phase val: 100%|██████████| 49/49 [00:26<00:00, 1.85it/s]

Loss: 0.00158
Accuracy: 0.97810
CharErrorRate: 0.00033

LR: 0.001

Epoch 011

Phase train: 100%|██████████| 449/449 [09:24<00:00, 1.26s/it]

Loss: 0.00469
Accuracy: 0.98964
CharErrorRate: 0.00127

Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.92it/s]
Loss: 0.00148
Accuracy: 0.97815
CharErrorRate: 0.00031

LR: 0.001

Epoch 012

Phase train: 100%|██████████| 449/449 [09:15<00:00, 1.24s/it]
Loss: 0.00493
Accuracy: 0.98926
CharErrorRate: 0.00132

Phase val: 100%|██████████| 49/49 [00:24<00:00, 2.00it/s]
Loss: 0.00145
Accuracy: 0.97790
CharErrorRate: 0.00034

LR: 0.001

Epoch 013

Phase train: 100%|██████████| 449/449 [08:55<00:00, 1.19s/it]
Loss: 0.00356
Accuracy: 0.99198
CharErrorRate: 0.00091

Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.92it/s]
Loss: 0.00096
Accuracy: 0.97890
CharErrorRate: 0.00020

LR: 0.0001

Epoch 014

Phase train: 100%|██████████| 449/449 [09:02<00:00, 1.21s/it]
Loss: 0.00415
Accuracy: 0.99138
CharErrorRate: 0.00106

Phase val: 100%|██████████| 49/49 [00:24<00:00, 2.02it/s]
Loss: 0.00099
Accuracy: 0.97900
CharErrorRate: 0.00019

LR: 0.0001


Epoch 015

Phase train: 100%|██████████| 449/449 [08:49<00:00, 1.18s/it]
Loss: 0.00269
Accuracy: 0.99328
CharErrorRate: 0.00071

Phase val: 100%|██████████| 49/49 [00:25<00:00, 1.90it/s]
Loss: 0.00091
Accuracy: 0.97910
CharErrorRate: 0.00017

LR: 0.0001

```
In [25]: plt.figure(figsize=(9, 5))
sns.lineplot(x=np.arange(1, n_epochs+1), y=train_losses, label='train loss')
ax = sns.lineplot(x=np.arange(1, n_epochs+1), y=val_losses, label='val loss')
ax.set(xlabel='Epoch', ylabel='Loss');
```



```
In [17]: torch.save(ocr.state_dict(), './models/fcnn_lstm_ctcloss_madgrad.pt')
```

Метрики на тестовой выборке

```
In [19]: ocr.score_test(test, metrics)
```

```
Phase test: 100%|██████████| 25/25 [00:13<00:00, 1.88it/s]
Loss: 0.00583
Accuracy: 0.99330
CharErrorRate: 0.00120
```

```
Out[19]: (0.005832448784559623,
{'Accuracy': 0.9932993299329933, 'CharErrorRate': tensor(0.0012)})
```

Error Handling

```
In [18]: outp = ocr.get_test_scores(test_dataset, torchmetrics.CharErrorRate(), device='cuda')
```

```
100%|██████████| 9999/9999 [01:09<00:00, 144.36it/s]
```

```
In [58]: scores = np.array(outp)
          scores_sorted_idx = scores.argsort()[::-1]
          sorted_scores = scores[scores_sorted_idx]
```

```
In [59]: sorted_scores[:70]
```

[illegible]

```
In [73]: @torch.no_grad()
def plot(start=0):
    plt.figure(figsize=(20, 10))
    for i in range(25):
        ax = plt.subplot(5, 5, i+1)
        img, label = test_dataset[scores_sorted_idx[i+start]]
        _, predicted = ocr.predict(img.cuda())
        ax.set_title(f'{label}-{predicted[0]}; {sorted_scores[i+start]:.3f}')
    plt.imshow(img.permute(1,2,0))
```


Топ 25 CER

```
In [71]: plot()
```



Топ 25-50 CER

In [75]: plot(25)



Отчет

Реализовано

- Свой класс данных
- Механизм аугментации изображений
- Предложенная в задании архитектура (FCNN + Bidirectional LSTM)
- Свой цикл обучения модели
- Энкодер/декодер символов, необходимый для данной архитектуры

Анализ ошибок модели

Было замечено, что модель чаще всего ошибается на картинках с некоторыми помехами (снег или потертости). Для решения проблемы были использованы аугментации из библиотеки `albumentations`:

- `PixelDropout` - меняет цвет случайных пикселей
- `CoarseDropout` - добавляет прямоугольники в случайных местах

Также использовался `ColorJitter` для лучшего предсказания на темных фото и на фото, с искаженным цветом.

Данные аугментации улучшили метрики, но не решили проблему полностью, модель все еще может ошибаться в некоторых

случаях.

Чаще всего она неверно предсказывает лишь 1 символ. Нередко это происходит в тех случаях, когда из-за шума можно спутать символы. Например, O-Q и F-E.

Наихудший CER модель имеет на 2 изображениях с искаженной перспективой и на изображениях с большим шумом.

RandomPerspective конкретно для этих 2 случаев не помог.

Возможные пути устранения:

- На последних эпохах заменять часть батча на "плохие" изображения
- Реализовать аугментации, имитирующие типичные помехи на изображениях.

Например, внизу часто заметен снег (или что-то еще), поэтому можно закрашивать нижние пиксели в форме синусоиды.

- Найти больше данных

Однако все равно будут случаи, когда однозначно нельзя определить символы

Метрики на тестовой выборке

Loss: 0.00583

Accuracy: 0.99330

CharErrorRate: 0.00120